

ARCHER Scripts

Tomas Lazauskas
t.lazauskas@ucl.ac.uk
University College London
lazauskas.net

September 25, 2016

Abstract

The help file of the set of scripts to execute jobs on ARCHER
(UK's National Supercomputing Service)

1 JobArray.bash

A simple job array script to execute many fhi-aims calculations using the same settings (i.e. the same control.in file).

An example is provided in the examples directory (Examples/JobArray).

1.1 Requirements

- input (directory): A directory where the system files and control files have to be saved.
- input/control.in (file): A standard control.in FHI-aims input file where the basis set and the computational parameters are set. The control file must be saved in the input subdirectory
- input/*.xyz (files): A set of xyz type structure files which will be optimised/evaluated with FHI-aims. At the moment the script only works with nanocluster type systems. If you would like to use it for periodic systems, please do not hesitate to contact me and I will make the changes.

- JobArray.bash (file) The main job array script which will pre-process the xyz files, prepare the directories and execute FHI-aims simulations. The results will be saved in the output directory.

1.2 Preparation

- Preparing directories and files:
Copy the main JobArray.bash file and create a new directory "input" in the dedicated directory on Archer. Save all the structure that you would like to optimise/evaluate with FHI-aims in the new input directory. Save the control.in file in the input directory.
- Setting the Archer job options:
Modify the JobArray.bash accordingly:
 - number of nodes (line 5)
e.g. #PBS -l select=5
will request 5 nodes (120 cores)
 - Project code (line 9)
e.g. #PBS -A projectX
the project code against which the job will be charged is set as "projectX"
 - number of cores per FHI-aims simulation (line 80)
e.g. aprun -n 24 aims > FHIaims\${i}.out &
each fhi-aims simulation will use 24 cores (1 node) and 5 simulations will be run at same time. It is important to make sure that the total requested number of cores is divisible by the number of cores per simulation to ensure the efficient use of the computational resources.

1.3 Execution

Simply by qsub'ing the JobArray.bash file